

# Gossip and Push-Sum Protocols: Convergence Time Analysis

## Team Members

- Samarth Vinayaka (UFID - 99665198)
- Pulkit Garg (UFID - 31125456)

## Overview

This project implements and analyzes the convergence time of the **Gossip** and **Push-Sum** algorithms on various network topologies using the Gleam language.

## Experimental Setup

- **Algorithms:** Gossip, Push-Sum
- **Topologies:** Full Network, Line, 3D Grid, Imperfect 3D
- **Network Sizes:** Multiple (e.g., 60, 80, 100, 120, 140 nodes)
- **Platform:** Gleam/Erlang on BEAM VM

## Gleam Packages Used

The following Gleam packages are used in this project:

- `gleam_stdlib`: Core standard library for Gleam, providing essential data structures and functions.
- `gleam_otp`: Provides OTP (Open Telecom Platform) abstractions for building concurrent and fault-tolerant applications.
- `gleam_erlang`: Enables interoperability with Erlang and access to BEAM-specific features.
- `gleeunit`: (dev dependency) Used for writing and running tests in Gleam.

## Project Modules Overview

- `project2_gossip.gleam`: Main entry point for running the project and orchestrating the simulation.
- `boss.gleam`: Coordinates the actors and manages the execution of the protocols.
- `node.gleam`: Defines the behavior and state of individual nodes in the network.
- `topology.gleam`: Handles the creation and management of different network topologies.
- `simple_gossip.gleam`: Contains a simplified or helper implementation of the gossip protocol.
- `time_util.gleam`: Provides utility functions for timing and measuring performance.
- `argv.gleam`: Handles parsing of command-line arguments for the application.


## What is Working

- Both Gossip and Push-Sum algorithms are fully implemented and can be run on all four topologies (Full Network, Line, 3D Grid, Imperfect 3D).
- The system successfully simulates large-scale networks, with Push-Sum handling up to 10,000 nodes and Gossip up to 3,000 nodes in some topologies.
- Convergence time is measured and reported for each run, and results are consistent with expectations for distributed protocols.
- The codebase is modular, allowing easy extension or modification of algorithms and topologies.

## Results: Convergence Time vs. Network Size


For each algorithm, we plot the convergence time as a function of network size for all four topologies. Each graph overlays the four topologies for direct comparison. We collected 40 result sets and presented aggregate findings.

### Push-Sum Algorithm



- The Push-Sum algorithm demonstrates a clear trend of increasing convergence time as the network size grows, but the rate of increase varies significantly across different topologies.
- The line topology consistently achieves the fastest convergence, while the 3D grid shows noticeably higher convergence times, especially at larger network sizes.
- The full network topology performs well for smaller sizes but does not scale as efficiently as the line topology, highlighting the impact of network structure on distributed algorithm performance.

## Gossip Algorithm



- The convergence time for the Gossip algorithm rises sharply as the network size increases, indicating that scalability is a significant challenge for this protocol.
- Among all topologies, the line topology consistently achieves the fastest convergence, while the 3D grid and imperfect 3D topologies show much slower convergence, especially at larger network sizes.
- The full network topology starts with relatively low convergence times for small networks but hits a threshold cap much earlier at higher node count as compared to other topologies.

## Largest Network Size Managed

The largest network size successfully managed for each topology and algorithm is as follows:

Algorithm	Full Network	Line Topology	3D Grid	Imperfect 3D
Gossip	2500	3000	3000	3000
Push-Sum	10000	10000	10000	10000

*Note: For Push-Sum, the implementation was able to handle up to 10000 nodes, while for Gossip the maximum tested was 3000 nodes.*

## Key Findings

- The network structure really matters. As the number of nodes grows, convergence slows down, but the effect is different depending on the topology.
- The line topology surprised me by being the most consistent and scaling the best overall, while the full network wasn't as efficient as I expected when the network got bigger.
- The 3D and imperfect 3D topologies worked, but they tended to slow things down compared to line or full, especially as the system grew.
- Between the two algorithms, Push-Sum clearly converges faster and can handle much larger networks than Gossip.
- The data also suggests that Push-Sum sometimes follows a sub-linear growth pattern, which is why it scales much better in larger setups compared to Gossip.

## Usage

```
# Build the project
gleam build
```

```
# Run gossip algorithm
gleam run -m project2_gossip <nodes> <topology> gossip
```

```
# Run push-sum algorithm
gleam run -m project2_gossip <nodes> <topology> push-sum
```

```
# Example:
gleam run -m project2_gossip 100 full gossip
gleam run -m project2_gossip 60 line push-sum
```

- <nodes>: Number of nodes (e.g., 60, 80, 100, 120, 140)
- <topology>: full, line, 3D, imp3D
- <algorithm>: gossip, push-sum